

Measured end-to-end latency: SDPA vs Python M1 vs M4 Triton (RTX 4060 Ti)

